



Modeling Convex Optimization Problems

CVX and CVXOPT

Vishal Gupta
Jan 31, 2013

Outline

- CVX Basics
 - What is CVX?
 - Convexity and DCP Convexity
- Advanced CVX
 - Dual variables
 - SDPs, GPs and MICPs
 - Solver settings
- CVXPY and CVX_OPT
 - CVXPY (brief)
 - Modeling language vs. solver
 - CVXOPT Basic Usage and documentation
 - Specializing Linear Algebra (time permitting)
- Course Wrap-up

Full Disclosure

- I ***strongly dislike*** Matlab
 - Hangs/becomes unstable
 - Many Coding conventions are frustrating (namespaces?)
 - Third-party libraries (when they exist) can be buggy
 - Expensive, commercial software
- *Why am I teaching CVX (a Matlab based software)?*

In my experience, CVX is the fastest way to prototype a convex optimization problem.

Outline

- CVX Basics
 - What is CVX?
 - Convexity and DCP Convexity
- Advanced CVX
 - Dual variables
 - SDPs, GPs and MICPs
 - Solver settings
- CVXPY and CVX_OPT
 - CVXPY (brief)
 - Modeling language vs. solver
 - CVXOPT Basic Usage and documentation
 - Specializing Linear Algebra (time permitting)
- Course Wrap-up

What is CVX?

- CVX is a Matlab based language for modeling convex optimization problems
 - Originally academic, now commercial (<http://cvxr.com/>)
 - Very intuitive, well-documented
 - Variety of back-end solvers (SeDumi, SDPT3, MOSEK, GUROBI)
- Well suited for:
 - Exploring small / medium problems
 - Comparing formulations
 - Prototyping optimizations in research
 - Coding small MILPs for Gurobi in MATLAB

What is CVX NOT?

- CVX is NOT
 - A solver
 - Tool to test convexity
 - Particularly clear with error messages (as of writing)
- (In my opinion), CVX is not well-suited for:
 - Limited support for mixed-integer convex problems... (Not yet mature)
 - Solving very large instances
 - Solving many, many small instances (as in simulation)

Learn by doing...

- The best way to learn CVX by examples
 - Second best: Consult documentation
 - <http://cvxr.com/cvx/doc/>
- Please open “testCVX.m” and run it.
- Questions:
 - What is this optimization doing?
 - What is the dimension of x ? The value of x_1 ?
 - How many non-zero values of x are there?

CVX Basics

- Problem surrounded by `cvx_begin ... cvx_end`
 - `cvx_begin quiet ... cvx_end` will suppress output
- Variables are declared first
 - Variable $x(n)$ ~ indicates a vector
 - Variable $y(m, n)$ ~ an n by m matrix
 - Variables $x(n)$ $y(m, n)$ z ~ a vector, a matrix and a scalar!
 - We'll see other possibilities later
- After `cvx_end`, automatically solves...
 - `cvx_optval` contains optimal value
 - Variables are replaced with their optimal values

Your turn...

The previous least-squares problem was badly conditioned.

- Modify the optimization to solve a regularized problem

$$\min \|\mathbf{r}\|_2 + \lambda \|\mathbf{x}\|_1$$

for $\lambda = 1$

- Embed this optimization into a script to solve the above for $\lambda = 2^{-3}, 2^{-2.5}, \dots, 2^{-3}$. Plot the trade-off curve. (Hint: Use a “for” loop. Suppress output in the loop.)

- Challenge Questions:

- Is CVX white-space sensitive?
- How would you incorporate the constraints

$$-10 \leq x_i \leq 10 \quad i = 1, \dots, n$$

Summary of the basics

- Wrap optimization in `cvx_begin`, `cvx_end`
- Write functions and constraints as you would write them in Matlab
- Automatically attempts to solve, and replaces variables with their optimal values
- Can embed within Matlab scripts

Review: Convexity

- Recall the definition of convexity:

$$f(\lambda \mathbf{x} + (1 - \lambda)\mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda)f(\mathbf{y}) \quad \forall \lambda \in [0, 1]$$


- Intuitively, convex functions are like bowls (not hats)

- Convex functions

- $\mathbf{a}^T \mathbf{x} + b$
- e^x, e^{-x}
- $\|\mathbf{x}\|$
- $\mathbf{x}^T \mathbf{Q} \mathbf{x}, \quad \mathbf{Q} \succeq \mathbf{0}$

- Non-convex functions

- $\mathbf{x}^T \mathbf{Q} \mathbf{x}, \quad \mathbf{Q}$ indefinite
- $\log x, \sqrt{x}$



Review: Convexity II

- Combining convex functions in certain ways preserves their convexity:
 - $f(x)$ convex implies $g(x) = f(Ax + b)$ is convex.
 - $f(x)$ concave implies $g(x) = -f(x)$ is convex
 - $f_1(x), f_2(x)$ convex implies $g(x) = c_1 f_1(x) + c_2 f_2(x)$ is convex if $c_1, c_2 > 0$
 - See *Convex Optimization* by Boyd and Vandenberghe for advanced examples

Convex Sets


- Recall the definition of a convex set

$$\mathbf{x} \in S, \mathbf{y} \in S \Rightarrow \lambda \mathbf{x} + (1 - \lambda) \mathbf{y} \in S$$

- Intuitively, convex sets are “round”

- Convex sets:

- $\{\mathbf{x} : \mathbf{A}\mathbf{x} = \mathbf{b}\}$
- \mathbb{R}_+^n
- $\{\mathbf{X} : \mathbf{X} \succeq \mathbf{0}\}$
- $\{(\mathbf{x}, t) : \|\mathbf{x}\| \leq t\}$



Images taken from Wikipedia: http://en.wikipedia.org/wiki/Convex_set

Convex Optimization Problems

- A minimization problem with a convex objective function and convex feasible region is called a ***convex optimization***
- Otherwise, we say it is ***non-convex***

Disciplined Convex Programming (DCP)

- **CVX does NOT model convex optimization**
 - Popular misconception
 - CVX models a subset of convex optimizations
- Specifically, CVX models problems obeying Disciplined Convex Programming (Grant, Boyd and Ye)

Disciplined convex programming imposes a set of convention or rules... Problems which adhere to the ruleset can [automatically] converted to solvable forms. Problem which violate the ruleset [even if they are convex] are rejected.

Excerpted from <http://cvxr.com/cvx/doc/intro.html>

DCP... intuitively

- There are certain *atomic* functions that CVX knows are convex
 - $e^x, \mathbf{a}^T \mathbf{x} + b, x^4, \|\mathbf{x}\|, \mathbf{x}^T \mathbf{Q} \mathbf{x} \ (\mathbf{Q} \succeq \mathbf{0})$
 - See [documentation](#) for full list
- CVX also knows a set of rules for combining convex functions to build new convex functions (DCP Rule set)
 - $\|\mathbf{A}\mathbf{x} + \mathbf{b}\|$ (composition of convex and affine function)
 - See [documentation](#) for full list

DCP (continued)

- We will say that a function is “DCP-convex” if it is either atomic and convex, or can be constructed from the atomic functions by the DCP rule set
- Similar definitions apply to “DCP-concave”, “DCP-affine”, “DCP-non-decreasing”, “DCP-non-increasing”, etc.

Valid optimization problems

- A Valid DCP objective is either
 - $\min f(\mathbf{x})$ where f is DCP-convex
 - $\max f(\mathbf{x})$ where f is DCP concave
- Valid DCP constraints are of the form
 - $f(\mathbf{x}) \leq g(\mathbf{x})$ where f is DCP convex, g is DCP concave
 - $f(\mathbf{x}) = g(\mathbf{x})$ where f, g are DCP-affine
 - $x \in S$ where S is an atomic, convex set (more on this later)
- Not all convex optimizations are a valid DCP
- Almost all can be transformed to be a valid DCP

Examples:

- Constraint: $\|\mathbf{x}\|_2 \leq 1$
 - $\|\mathbf{x}\|_2$ is convex and DCP convex, 1 is DCP concave
 - This is a **valid constraint**
- Constraint: $\|\mathbf{x}\|_2 \geq 1$
 - $\|\mathbf{x}\|_2$ is not DCP concave. In fact, it's DCP convex (and convex).
 - This is an **invalid constraint**.
- Constraint: $\sqrt{x_1^2 + x_2^2} \leq 1$
 - $\sqrt{x_1^2 + x_2^2}$ is convex, but NOT DCP convex
 - Hence, this is an **invalid constraint**
 - Can be rewritten in DCP way

Error using cvx/sqrt (line 61)
Disciplined convex programming
error:
Illegal operation: sqrt({convex})

Some practical advice....

- It might seem that you have to memorize all the DCP rules in order to know what you can and cannot model.
- Personally, I have not memorized the rules.
 - In my experience, most convex functions you come across in practice are DCP convex.
 - When they aren't, scanning the documentation for CVX for two minutes usually suggests the correct reformulation.
 - "Easier to ask forgiveness than to get permission."

Let's try some examples!

Your turn...Maximal entropy problems

Let X be random variable distributed on $\{1, 2, \dots, 10\}$ such that

$$\mathbb{P}(X = i) = p_i \quad i = 1, \dots, 10$$

The entropy of X , denoted $H(X)$ is defined by:
(Entropy is a concave function of the p 's).

$$H(X) = - \sum_{i=1}^{10} p_i \log p_i$$

Suppose we know that $\mathbb{E}[X] = 4$

Formulate an optimization problem to find the maximal entropy distribution with these properties.

Max Entropy II

$$\begin{aligned} \max_{\mathbf{p}} \quad & - \sum_{i=1}^{10} p_i \log(p_i) \\ \text{s.t.} \quad & \sum_{i=1}^{10} p_i = 1, \quad \mathbf{p} \geq \mathbf{0} \\ & \sum_{i=1}^{10} i p_i = 4 \end{aligned}$$

1. Solve this problem in CVX.
2. Add the constraint that the standard deviation of X is at most 4. Resolve.
3. Read the output of the solver. Notice anything different?

Summary on DCP Convexity

- CVX only models DCP-convex problems
- Remember my advice
 - “Easier to ask forgiveness than permission”
- When in doubt, the documentation is very good:
 - <http://cvxr.com/cvx/doc/index.html>

Outline

- CVX Basics
 - What is CVX?
 - Convexity and DCP Convexity
- Advanced CVX
 - Dual variables
 - SDPs, GPs and MICP
 - Solver settings
- CVXPY and CVX_OPT
 - CVXPY (brief)
 - Modeling language vs. solver
 - CVXOPT Basic Usage and documentation
 - Specializing Linear Algebra (time permitting)
- Course Wrap-up

Dual Variables

- Possible to extract the dual variables for convex models.
 - Frequently no extra cost computational cost
 - Always non-negative for inequality constraints
- CVX automatically determines the dimensions of the dual
 - Error if you try to specify
 - Useful for matrix inequalities like $x \leq 1$
- After solving, dual variables (like decision variables) populated with optimal values

Example 1 Constraint:

- Constraint our Max-Entropy Problem (maxEntropy.m)

```
%Constrain mean  
supp * p == 4;
```

- With dual variable (See SimpleDual.m)

```
%Incorporate Dual variable  
dual variable lambda  
lambda : supp * p == 4;
```

Size of lambda implicitly
determined to be 1

- Specifying dimension yields error

```
dual variable lambda(1)
```

Invalid dual variable
specification: lambda(1)

Example: Vector Constraint

- Constraint from our Regularized Least Squares

$$r == y - A * x$$

- With dual variable:

```
dual variable lambda  
lambda: r == y - A * x
```

Size of lambda implicitly
determined to be $m = 100$

Constraints in a Loop I

- Consider following optimization (LoopedLeastSquares.m):

```
cvx_begin
    variables x(n) r(m)
    for i = 1:m
        r(i) == y(i) - A(i, :)*x
    end
    minimize norm(r, 2) + norm(x, 1)
cvx_end
```

- How do we incorporate dual variables now?

Constraints in a Loop II

- Use Matlab's cell array feature

```
cvx_begin
    variables x(n) r(m)
    dual variable p{m}
    for i = 1:m
        p{i} : r(i) == y(i) - A(i, :) * x
    end

    minimize norm(r, 2) + norm(x, 1)
cvx_end
```

- Simultaneously creates (m) dual variables
 - Notice the {} instead of the usual () notation
 - Only really necessary for looping constructs

Your turn ... Max Entropy revisited

- Open the file “MomentProblems.m “
 - What is this code doing?
 - Is it possible to rewrite this optimization without the loop?
- Add dual variables to this optimization for:
 - The normalization constraint
 - Each of the moment constraints
 - Plot the dual variables for each moment

Max Entropy Solution

- See “MomentProblemsWithDual.m”

SDP mode

- You should have read a little bit about SDPs as part of your homework.

$$\begin{aligned} \min \quad & 3x_{11} - 2x_{23} \\ \text{s.t.} \quad & 5x_{11} + 2x_{23} - x_{33} = 1 \\ & 2x_{22} + 3x_{33} \leq 2 \\ & \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{12} & x_{22} & x_{23} \\ x_{13} & x_{23} & x_{33} \end{pmatrix} \succeq \mathbf{0} \end{aligned}$$

- Recall:
 - Optimization over matrices
 - Objective and constraints are linear functions of the entries
 - Matrix is constrained to be symmetric and positive semidefinite

Take my word for it...

- SDPs are extremely common and very useful...
 - Applications to probability, control theory, combinatorial optimization
 - Multiple courses at MIT that cover their usage
- For today, just take my word that if you work long enough in optimization, eventually you will want to solve an SDP

Your first SDP in CVX

- (See SimpleSDP.m)

$$\begin{array}{ll}\min & 3x_{11} - 2x_{23} \\ \text{s.t.} & 5x_{11} + 2x_{23} - x_{33} = 1 \\ & 2x_{22} + 3x_{33} \leq 2 \\ & \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{12} & x_{22} & x_{23} \\ x_{13} & x_{23} & x_{33} \end{pmatrix} \succeq \mathbf{0}\end{array}$$

```
cvx_begin
    variable X(3, 3)

    5 * X(1,1) + 2 * X(2, 3) - X(3, 3) == 1
    2 * X(2, 2) + 3 * X(3, 3) <= 2
    X == semidefinite(3)

    minimize 3 * X(1, 1) - 2 * X(2, 3)
cvx_end
```

Other matrix operations

- $5x_{11} + 2x_{23} - x_{33} = 1 \Leftrightarrow \mathbf{C} \circ \mathbf{X} == 1, \quad \mathbf{C} = \begin{pmatrix} 5 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & -1 \end{pmatrix}$
 - `trace(C * X) == 1`
 - `vec(C)' * vec(X) == 1` (my preference... why?)
- $x_{11} + x_{22} + x_{33} = 1 \Leftrightarrow \text{tr}(\mathbf{X}) = 2$
 - `trace(X) == 2`

Other Matrix Operations Ct'd

- $\begin{pmatrix} 1 & y^T \\ y^T & X \end{pmatrix} \succeq 0$ (X and y are variables)

```
variables y(3, 1) X(3, 3)
variable Z(4, 4)
Z == semidefinite(4)
Z(1, 1) == 1
Z(2:end, 1) == y
Z(2:end, 2:end) == X
```

```
variables y(3, 1) X(3, 3)
[ 1, y';
  y, X ] == semidefinite(4)
```

Preferred method... why?

Your turn... Asset Correlation

- Consider three types of commodities futures: Oil, Natural Gas, and Electricity. (Index them 1, 2, 3)

- Suppose we know from the markets that

$$\begin{aligned} .6 &\leq \rho_{12} \leq .8 \\ .95 &\leq \rho_{23} \leq .98 \end{aligned}$$

- And assume we know that

$$\sigma_1 = .6, \quad \sigma_2 = .4, \quad \sigma_3 = 1.2$$

- What is the value of ρ_{13} ?

Formulate as an SDP

- $$\begin{array}{ll} \min & \rho_{13} \\ \text{s.t.} & .6 \leq \rho_{12} \leq .8 \\ & .95 \leq \rho_{23} \leq .98 \\ & \begin{pmatrix} .6^2 & \rho_{12} * .6 * .4 & \rho_{13} * .6 * 1.2 \\ \rho_{12} * .6 * .4 & .4^2 & \rho_{23} * .4 * 1.2 \\ \rho_{13} * .6 * 1.2 & \rho_{23} * .4 * 1.2 & 1.2^2 \end{pmatrix} \succeq \mathbf{0} \end{array}$$

(Other constraints here...)

- Code this up in CVX. What is optimal solution?
- Change your code to instead minimize

$$\min e^{\rho_{13}} + \|\rho_{12} - \rho_{23}\|$$

SDP Review

- Constraining a matrix to be PSD is a DCP-Convex constraint
 - There are other structured classes which are also DCP Convex (e.g. symmetric, hermitian, toplitz)
- We can combine SDP constraints with other DCP-convex constraints and objectives
- There is a separate “SDP mode” for CVX
 - Useful if you are coding up a large set of LMIs
 - Consult the documentation. Very straightforward.

Geometric programming

- GPs are a more advanced class of optimization problems
 - By a suitable transformation, they can be made DCP convex
- Check the documentation for details (GP mode)
 - [GP tutorial](#)
 - [CVX documentation on GPs](#)

Mixed Integer Programming

- It is possible to solve mixed integer convex problems
 - This is still a developing field
 - With non-Gurobi solvers, uses a proprietary branch and bound
 - Need to be careful in terms of computational power etc.
- For MILP, MIQP possible to use Gurobi as back-end
 - Requires a “professional” license (free for academics)
 - Set up is relatively easy. See the documentation.
- To use: Simply define variables with integer type

```
variable p(10) integer  
variable z(10) binary
```

Fine-tuning Solver Settings

- Change the precision
 - `old_precision = cvx_precision('high')`
 - Possible choices: low, medium, default, high, best
- Change the solver
 - `cvx_solver sedumi`
 - Possible choices: sedumi, sdpt3, gurobi, mosek
 - Last two require separate installation and a (free) academic license
- Pass specific settings to solver
 - Badly documented
 - `cvx_solver_settings("method", 1)` // Dual simplex alg for Gurobi
 - Key-value pairs are solver dependent. Will throw if not supported.

Summary of Advanced Features

- Dual Variables
 - Size is implicitly determined
 - Use cell array {} for loops
- Semidefinite Programming
 - Membership to the semidefinite cone is DCP-convex
 - Many ways to write same constraints.
 - SDP mode may be useful to you
- Fine tuning solver settings
 - Generally shouldn't need to alter these
 - Documentation is a little spotty. Guess and check is best bet.

Where do I learn more?

- CVX Documentation is your friend...
 - <http://cvxr.com/cvx/doc/index.html>
- There are a ton of examples that ship with the distribution
 - <path to CVX>/examples
 - Many of them parallell the textbook *Convex Optimization* by Boyd and Vanderberghe
- Experiment on your own!
 - Most things in CVX are intuitive.

Outline

- CVX Basics
 - What is CVX?
 - Convexity and DCP Convexity
- Advanced CVX
 - Dual variables
 - MIPs, SDPs, GPs
 - Solver settings
- CVXPY and CVX_OPT
 - CVXPY (brief)
 - Modeling language vs. solver
 - CVXOPT Basic Usage and documentation
 - Specializing Linear Algebra (time permitting)
- Course Wrap-up

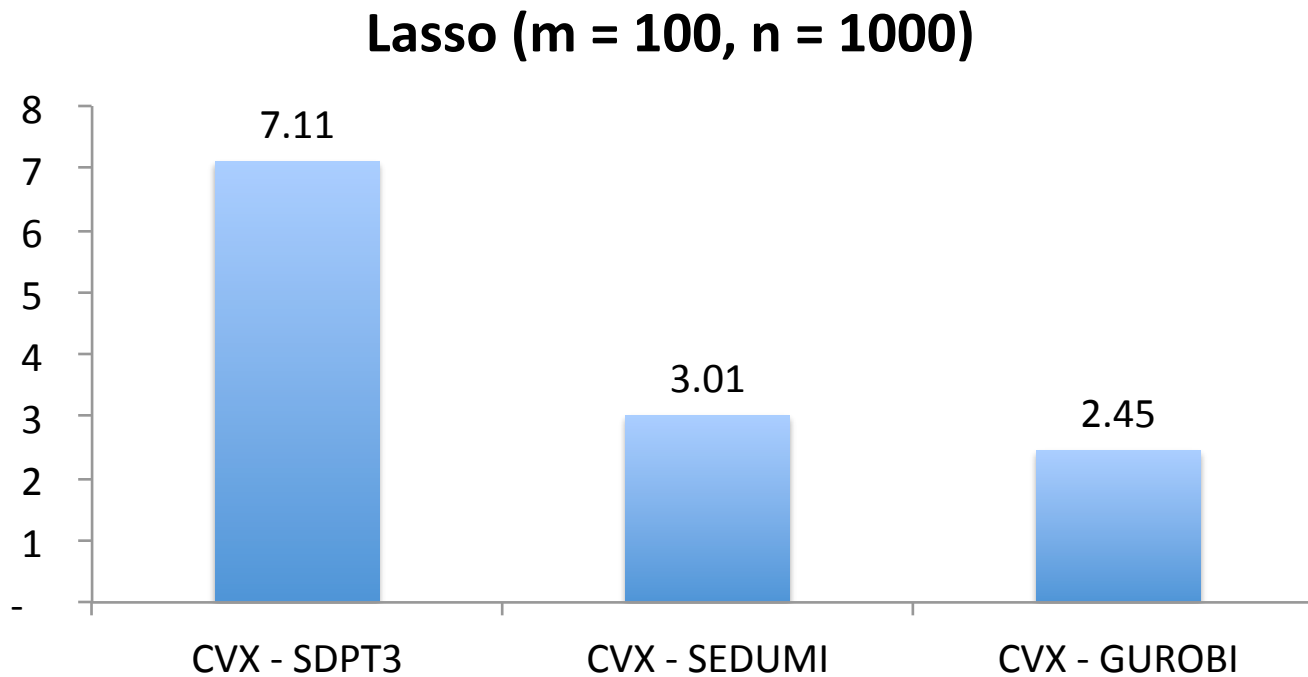
Performance Bottlenecks... When not to use CVX?

- Recall our friend Lasso Regression

$$\min_{\mathbf{x}} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2 + \lambda \|\mathbf{x}\|_2$$

- Dim(A) = 100 by 1000, dense
 - $\lambda \in \{2^{-5}, 2^{-4.75}, \dots, 2^5\}$
 - Sparse regression contexts
-
- Solving this problem in CVX takes about 7s per instance
-
- Is this slow? Why?

Could vary the solver...



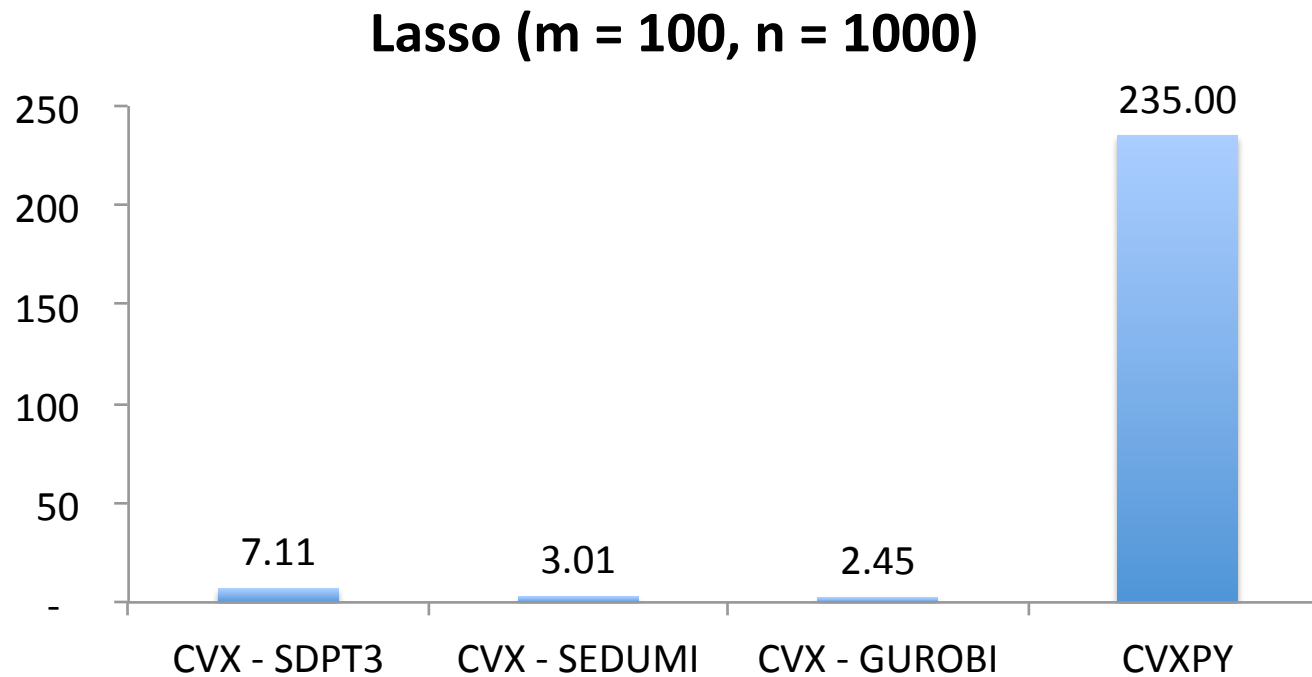
Helps a little... Can we do better?

CVXPY

- Similar to CVX but implemented in Python
- Once you understand CVX, simple to learn CVXPY from documentation
- Some sample code (“cvxpy_timing.py”)

```
x = variable(n, 1, name="x")
obj = minimize( norm2( y - A * x) + norm1(x))
cnsts = []
p = program(obj, cnsts)
p.solve()
```


And the results...



What happened?

An important warning

- Slow code is slow no matter what language it is in...
- CVXPY was developed as academic software and then seemingly abandoned.
- CVX on the other hand is now sold commercially.
 - Has undergone multiple releases
 - At this point is highly optimized as a modeling language

From modeling language to solver

- A smarter approach: bypass the modeling language and go straight to the solver
- CVXOPT is a collection of solvers for conic programs,
 - LP, QP, GP and generic cone solvers
 - Sparse Matrix functionality, and high-performance linear algebra
- Challenge is that solvers expect you to specify your problem in **standard form**
 - Converting your problem (on paper) can be tricky
 - Coding up the reformulation can be tedious


Example QP Solver

$$\begin{array}{ll}\min_{\mathbf{x}} & \frac{1}{2}\mathbf{x}^T \mathbf{P} \mathbf{x} + \mathbf{c}^T \mathbf{x} \\ \text{s.t.} & \mathbf{G} \mathbf{x} \leq \mathbf{h} \\ & \mathbf{A} \mathbf{x} = \mathbf{b}\end{array}$$

- This is the standard form cvxopt expects for QPs
- Solve by calling `cvxopt.qp(P, c, G, h, A, b)`
- Can frame our Lasso problem in this form?

$$\min_{\mathbf{x}} \|\mathbf{y} - \mathbf{A} \mathbf{x}\|_2 + \lambda \|\mathbf{x}\|_2$$

Board Work



Solution

$$\text{“}P\text{”} = \begin{pmatrix} 2A^T A & 0 \\ 0 & 0 \end{pmatrix}$$

$$\text{“}c\text{”} = \begin{pmatrix} -2A^T b \\ \lambda e \end{pmatrix}$$

$$\text{“}G\text{”} = \begin{pmatrix} I & -I \\ -I & -I \end{pmatrix}$$

$$\text{“}h\text{”} = \begin{pmatrix} 0 \\ \vdots \\ 0 \end{pmatrix}$$

$$\text{“}A\text{”} = \text{“}b\text{”} = 0$$

Coding this up in CVXOpt

- CVXOpt has *many* convenience for building up (sparse) matrices to pass to the optimizers

$$\text{matrix}(1, (2, 3)) = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

Repeat a value in given shape

$$\text{matrix}([[1, 2, 3], [4, 5, 6]]) = \begin{pmatrix} 1 & 4 \\ 2 & 5 \\ 3 & 6 \end{pmatrix}$$

List of Lists Populated by Column

$$\text{matrix}([[A, 1], [0, B]]) = \begin{pmatrix} A & 0 \\ 1 & B \end{pmatrix}$$

Works for Block Matrices, too

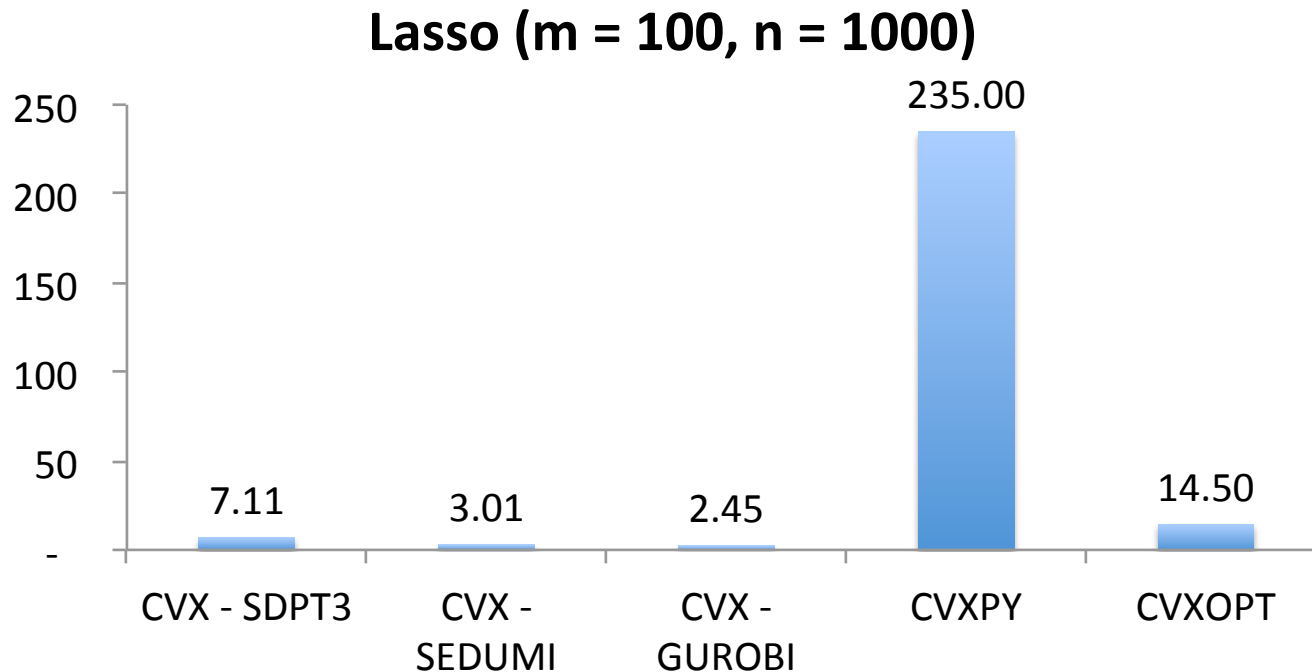
$$\text{spdiag}([A, B, 1]) = \begin{pmatrix} A & 0 & 0 \\ 0 & B & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

Can also create sparse block-diagonal matrices

Let's (together) code up our Lasso problem

- “cvxopt_timing.py”

And the results....



- Certainly better than CVXPY
- But the CVXOPT cone solvers are not quite as good as others

Customizing Linear Algebra

- How does Newton's method work? Where is the computational burden?
 - Interior point solvers have a similar problem. Need to take solve a large linear system from KKT conditions
- Often, if your problem has special structure, you can customize the linear algebra to solve this system
 - General consensus is that this is tedious, but the benefits are enormous
 - The paper by Andersen, et.al on website some examples and explains the background mathematics

Example...

- For the case of lasso regression, the KKT equations look like:

$$\begin{pmatrix} A^T A + W_1^{-2} + W_2^{-2} & W_2^{-2} - W_1^{-2} \\ W_2^{-2} - W_1^{-2} & W_2^{-2} + W_1^{-2} \end{pmatrix}$$

where W_1 , and W_2 are diagonal.


- How is this matrix structured?
- Turns out we can reduce this system to a much smaller system...

$$(A^T A + D)x = rhs$$

$$D = 4(W_1^2 + W_2^2)^{-1}$$

And the results...

- Take a look at “l1regls.py” to see how this is implemented...



Morale of the story...

- CVX is my choice for *rapid* prototyping
 - Especially good when you want to change the formulation repeatedly
- If you need to code in python
 - CVXPY is an acceptable choice for small problems
 - CVXOPT is a much better choice.
- If you have the need, patience and skill, customizing linear algebra is a good way to speed things up *significantly*.

Outline

- CVX Basics
 - What is CVX?
 - Convexity and DCP Convexity
 - Advanced CVX
 - Dual variables
 - SDPs, GPs and MICPs
 - Solver settings
 - CVXPY and CVX_OPT
 - CVXPY (brief)
 - Modeling language vs. solver
 - CVXOPT Basic Usage and documentation
 - Specializing Linear Algebra (time permitting)
- Course Wrap-up

Course Contents

- We've covered many tools in a short time...
 - [R] for Data Analysis
 - D3.js for Interactive Visualization
 - Python (PuLP) for Linear Optimization Modeling
 - Java and CPLEX for customizing linear optimization solvers
 - Distributed Computing on the EC2 Cloud
 - CVX and CVXOpt for Convex Optimization Modeling
- Why?
 - This suite of tools spans the various phases of research, development and implementation.
 - It's always important to think about what is the right tool for the job.

Situation 1

- You have a dataset of transaction data for a large retailer organized by customer. You're looking to build a new model to describe customer preferences.
- What would you use to do the initial data analysis?
- How would you implement and back-test various models to test their predictive power?

Situation 2

- You do research in linear integer programming. You're interested in developing heuristics and algorithms for a specific class of problems (say tripartite graph matching).
- To gain some intuition, you'd like to solve some sample instances, and study the structure of the solutions.
- What tools would you use to do this?

Situation 3

- After studying the solutions, you come up with (paper and pencil) a number of possible linear programming relaxations and would like to compare their relative strength.
- What tools would you use?

Situation 4

- Your genius and hardwork has finally paid off! You have a great new linear programming based heuristic.
- You would like to run the heuristic on many simulated instances to collect computational evidence for your forthcoming journal paper.
- How would you do this?
- How would you make a high-quality graphic for the paper?

Situation 5

- Even better than a paper, a company wants to use your heuristic as part of their operations.
- How would you demonstrate your work to non-technical audiences?



Thank you